

Default Question Block

Thank you for considering taking part in this study.

This survey is part of a Master's thesis project on leveraging Large Language Model to generate short narratives about regulatory concepts of the EU Artificial Intelligence Act (AI Act) that startup companies are likely to encounter or struggle with, according to scientific literature. The purpose of the generated narratives is to simplify the meaning of these regulatory concepts for non-experts and to clarify how this target group is expected comply with such concepts.

The answer of the research question will be based on an experiment with students from the University of Amsterdam as participants who qualify as non-experts regarding the AI Act. These participants will read both generated narratives and legal text summaries on the same regulatory concepts, after which their understanding is evaluated through survey questions.

For each of the four selected regulatory concepts for this study, you will read three different short narratives ranging from 200-400 words, that are based on three different types of prompt engineering, and one generated multiple-choice question, that is ultimately used for the evaluation of the participants' understanding. After this, you are asked to evaluate the generated narratives and the associated multiple-choice question through a ranking question and a select number of Likert-type

scale questions, after which there is an optional opportunity for you to provide further feedback.

This survey is estimated to take about 20-25 minutes and your responses will be used for academic research purposes only.

By continuing, you confirm that you have read this information and agree that your responses may be used to support and validate the planned main experiment.

Regulatory Concept #1:

Below, you can find a link that takes you to *article 9* from the AI Act about *Risk Management System*. To prevent interruption or loss of your progress in this survey, it is advised to right-click the hyperlink and to open the link on another tab. It is **not mandatory** to read this, but it is available if you would like to scan through the original legal article, for example, as this might help you answer the upcoming questions.

<https://artificialintelligenceact.eu/article/9/>

Additional information:

From this point onward, you will encounter a few terms throughout this survey that are related to specific Large Language Models (LLMs) prompting strategies. Reading this short introductory text is entirely optional and is also **not required** to complete the survey. It is provided solely to offer additional context, prevent potential confusion, and ensure clarity.

Zero-shot prompting: prompting LLMs with only one carefully crafted prompt, which excludes the need for extensive training data, follow-up prompts, or given examples as a reference.

Contextual prompting: prompting LLMs by adding detailed context, which includes all descriptive information regarding the request and the desired output.

Rephrase and Response prompting: prompting LLMs by asking the model to expand and rephrase your request, and to respond to the new request that is created by the model itself.

Please read the following three generated narratives. After this, you are asked to rank these narratives from worst to best, in terms of how well they represent compliance with Risk Management System, as described in article 9 (1 best - 3 worst). You can find this question by scrolling down, as it is positioned right below the three narratives:

Zero-shot prompting:

Sophie's team was building an AI tool to help hospitals prioritise urgent patient cases, and before anyone was allowed to use it, she gathered the engineers, doctors, nurses, and legal adviser around one table. Instead of asking, "Does the system work?", she asked, "Who could be harmed if it works badly, is misunderstood, or is used in a way we can already predict?" They listed risks to patient safety, unfair outcomes for certain groups, overreliance by busy staff, and mistakes that might appear only after the tool was used in real hospital conditions. For each risk, they decided what could be designed out of the system, what needed a

warning or control, and what doctors needed to know before using it. They tested the tool again and again, not only in ideal conditions but also with messy cases, tired users, and unusual patient files. Some risks could not be removed completely, so Sophie wrote down why the remaining risk was acceptable and what safeguards would stay in place. After launch, the team kept watching how the system behaved, collected incident data, reviewed the risk file regularly, and updated the tool whenever new problems appeared. To Sophie, compliance was not a form signed at the end; it was more like keeping a hospital fire plan alive: identify dangers early, reduce them as far as possible, test the precautions, explain the remaining risks clearly, and keep checking for new ones throughout the system's life.

Contextual prompting:

Mila's AI startup had built a tool that helped hospitals decide which patient messages looked urgent, but before launching it, her team treated the system like a new colleague who needed careful supervision, not like a finished product. They first sat around a table and asked what could go wrong if the tool was used exactly as intended: could it miss a serious symptom, make unfair suggestions for certain patient groups, or push nurses to trust it too much? Then they imagined realistic mistakes people might make, like using it for children even though it had only been tested on adults, or relying on it during very busy night shifts without checking the output properly. For every risk they found, they wrote down how serious it could be, how likely it was, and what they could do about it. Some problems were fixed in the design, such as adding warnings when the model was uncertain. Others could not be removed completely, so they added safeguards, like human review for urgent cases, clear instructions for hospital staff, and training notes explaining when the tool should not be used. Before release, they tested the system again and again, including with realistic hospital scenarios,

to see whether the safeguards actually worked. Even after the product went live, Mila kept a shared risk log open: if hospitals reported confusing outputs, unexpected behaviour, or new situations the team had not predicted, the team reviewed the risks again and updated the system or its instructions. By the time investors asked whether the company was “AI Act ready,” Mila did not point to a single checklist. She pointed to the living process her team had built: identify risks, judge them, reduce them where possible, explain the remaining ones, test before launch, and keep learning after launch. That was their risk management system.

Rephrase and Respond prompting:

Maya worked at a small company building an AI tool that helped hospitals decide which patients should be seen first in the emergency department. At first, the team was excited because the system seemed fast and accurate, but before anyone was allowed to use it, Maya gathered the doctors, nurses, engineers, and patient safety officer around one table and asked a simple question: “What could go wrong, even if everyone uses this exactly as intended?” They listed the obvious dangers first: the AI might underestimate a child’s symptoms, misunderstand rare conditions, perform worse for people whose data was underrepresented, or make staff trust its recommendation too much. Then Maya asked a second question: “How might people misuse it in a way we can reasonably predict?” The team imagined a crowded night shift where a tired nurse might follow the AI score without checking the patient, or a hospital manager might use the tool outside the emergency department even though it was not designed for that. Instead of treating these worries as paperwork, Maya turned them into design decisions. The system was changed so that it could not make the final decision by itself, displayed warnings when cases were uncertain, explained the main reasons behind its recommendation, and clearly

told staff when the tool should not be used. The team tested it again and again, including with difficult cases, children’s cases, and examples from different patient groups, until the remaining risks were reduced as much as technically possible and judged acceptable. Every important choice was written down: the risks found, the tests run, the trade-offs made, the safeguards added, and the instructions hospitals would need. When the tool finally went live, Maya did not close the risk file and move on. She set up a routine to review incident reports, user feedback, and performance data from real hospitals, because new risks could appear after deployment. Months later, when the team noticed that the AI was less reliable in a certain type of paediatric case, they updated the system, revised the instructions, and recorded the change.

Zero-shot prompting
Contextual prompting
Rephrase and Respond prompting

Based on the one generated narrative that you just ranked as the best representation of Article 9 - Risk Management System, please evaluate the following statements below:

	Strongly disagree	Somewhat disagree	Neutral	Somewhat agree	Strongly agree
The narrative contains correct information, avoiding any errors or inaccuracies compared to Art. 9	☐	☐	☐	☐	☐
The narrative is based on Art. 9, excluding any extraneous or unnecessary information	☐	☐	☐	☐	☐

	Strongly disagree	Somewhat disagree	Neutral	Somewhat agree	Strongly agree
The narrative contains relevant information and represents Art. 9 well	☐	☐	☐	☐	☐

How well does the following multiple-choice question examine whether non-experts have understood the regulatory concept of Risk Management System, as described in Article 9 of the AI Act?

Question: In the EU AI Act’s risk management system for high-risk AI systems, what is the main responsibility of the provider?

- (A) To prove before launch that the AI system is accurate enough, after which no further risk review is needed unless a serious incident occurs.
- (B)** To maintain a continuous process that identifies, evaluates, reduces, tests, documents, and regularly updates controls for risks to health, safety, and fundamental rights throughout the AI system’s lifecycle.
- (C) To transfer responsibility for all remaining risks to the human users by giving them warnings, training materials, and instructions before deployment.
- (D) To eliminate every possible risk completely before the AI system is placed on the market, even if doing so makes the system unable to function.

Not well ☒ at all	Slightly ☒ well	Moderately ☒ well	Very ☒ well	Extremely ☒ well
--	--	--	--	---

Do you have any further feedback on how your best ranked narrative and the multiple-choice question could be improved? (optional)



art 10

Regulatory Concept #2:

Below, you can find a link that takes you to *article 10* from the AI Act about *Data and Data Governance*. To prevent interruption or loss of your progress in this survey, it is advised to right-click the hyperlink and to open the link on another tab. It is **not mandatory** to read this, but it is available if you would like to scan through the original legal article, for example, as this might help you answer the upcoming questions.

<https://artificialintelligenceact.eu/article/10/>

Please read the following three generated narratives. After this, you are asked to rank these narratives from worst to best, in terms of how well they represent compliance with Data and Data Governance, as described in article 10 (1 best - 3 worst). You can find this question by scrolling down, as it is positioned right below the three narratives:

Zero-shot prompting:

Karen ran a small company that built an AI tool to help hospitals plan patient appointments, but before she let the system learn from old hospital records, she treated the data like ingredients for medicine. She

asked where every record came from, why it had first been collected, and whether it was allowed to be used for this purpose. Her team cleaned misspelled entries, checked labels, removed duplicates, and wrote down every assumption they were making, especially what each data point was supposed to represent. When they noticed that rural clinics and older patients were barely present in the training data, Karen stopped the launch, added better examples, and tested whether the system worked fairly for the people who would actually use it. A colleague found that one pattern could accidentally make appointments harder to get for patients with limited mobility, so they changed the dataset and added checks to catch that kind of bias early. Only when the training, validation, and testing data were relevant, representative, as complete and error-free as possible, and suited to the real hospital setting did Karen approve the system, keeping notes so anyone could later understand how the data had been chosen, prepared, checked, and improved.

Contextual prompting:

Mary and Jonas had just launched their Amsterdam AI startup, which built a tool to help hospitals decide which patients should be invited for an early diabetes screening. At first, their model looked impressive: it was fast, accurate in their demo, and investors loved it. But before they offered it to real clinics, Mary stopped the team and said, “We cannot just ask whether the model works. We need to ask whether the data deserves to be trusted.” So they opened a shared document and traced every dataset back to its source: which hospital it came from, why it had originally been collected, whether patients had been informed properly, and whether the data was meant for this kind of prediction. Then they checked how the data had been cleaned, labelled, updated, and combined. One intern noticed that some older records used different names for the same diagnosis, while another found missing values in

data from smaller clinics. Instead of ignoring those problems, the team wrote them down and decided how to fix or limit them. Jonas then compared the dataset with the people the system would actually be used for. He realised that most training data came from urban hospitals and did not include enough patients from rural areas, younger adults, or certain migrant communities. “If we train only on the people who are easiest to find,” he said, “the system may fail exactly when it meets people outside that pattern.” The team added more suitable data where they lawfully could, tested whether the model performed differently across groups, and looked for biases that could unfairly delay screening for some patients. When they needed sensitive information to check whether bias existed, they used it only where strictly necessary, protected access carefully, and deleted it once the bias check was done. By the time the startup met its first hospital client, Mary did not simply say, “Our AI is accurate.” She showed a record of the choices they had made: where the data came from, what it was supposed to represent, what gaps they found, how they reduced bias, and why the final training, validation, and testing data were suitable for the medical setting in which the tool would be used. The hospital’s compliance officer smiled and said, “So your data governance is not a folder you created at the end. It shaped the system before patients ever depended on it.”

Rephrase and Respond prompting:

At a city hospital, Mira was asked to help prepare a new AI tool that would support doctors in spotting urgent cases in the emergency department. Before anyone let the system near real patients, she treated the data like ingredients for a meal that many different people would have to eat safely. She did not just gather old patient files and hand them to the developers. First, she asked where every record came from, why it had originally been collected, and whether it was appropriate to use it for this new purpose. Then her team cleaned the files, corrected obvious

mistakes, checked the labels, removed duplicates, and wrote down the assumptions they were making, such as what certain symptoms or test results were supposed to represent. When the first model looked accurate, Mira still did not celebrate. She noticed that most of the training data came from large urban hospitals, while the tool would also be used in smaller regional clinics. So she added more suitable examples, checked whether older patients, women, and minority groups were properly represented, and tested whether the system made worse recommendations for any group. When the team found that one group was being flagged less accurately, they paused, investigated the cause, and adjusted the data and process before continuing. By the time the hospital approved the tool, Mira could explain not only that the AI had been trained, tested, and validated, but that its data had been chosen, prepared, checked for bias, and matched to the real setting where the system would be used.

Zero-shot prompting
Contextual prompting
Rephrase and Respond prompting

Based on the one generated narrative that you just ranked as the best representation of Article 10 - Data and Data Governance, please evaluate the following statements below:

	Strongly disagree	Somewhat disagree	Neutral	Somewhat agree	Strongly agree
The narrative contains correct information, avoiding any errors or inaccuracies compared to Art. 10	☐	☐	☐	☐	☐
The narrative is based on the Art. 10, excluding any extraneous or unnecessary information	☐	☐	☐	☐	☐
The narrative contains relevant information and represents Art. 10 well	☐	☐	☐	☐	☐

How well does the following multiple-choice question examine whether non-experts have understood the regulatory concept of Data and Data Governance, as described in Article 10 of the AI Act?

Question: A startup is preparing a high-risk AI system for use in hospitals. Which approach best reflects the data governance requirements described in Article 10?

- (A) The startup uses a large hospital dataset because the model performs well in a demo, but it does not investigate whether the data reflects the patients and clinical settings where the system will actually be used.
- (B)** The startup checks where the data came from, why it was collected, how it was cleaned and labelled, whether important groups or settings are missing, whether bias could affect outcomes, and whether the training, validation, and testing data are suitable for the system’s intended use.
- (C) The startup focuses mainly on whether the AI system is accurate overall, because Article 10 is about model performance rather than the way datasets are collected, prepared, and governed.
- (D) The startup avoids looking at sensitive group information under all

circumstances, even when this would be strictly necessary and carefully protected to detect and correct bias.

Not well at all

Slightly well

Moderately well

Very well

Extremely well

Do you have any further feedback on how your best ranked narrative and the multiple-choice question could be improved? (optional)

art 11

Regulatory Concept #3:

Below, you can find a link that takes you to *article 11* from the AI Act about *Technical Documentation*. To prevent interruption or loss of your progress in this survey, it is advised to right-click the hyperlink and to open the link on another tab. It is **not mandatory** to read this, but it is available if you would like to scan through the original legal article, for example, as this might help you answer the upcoming questions.

<https://artificialintelligenceact.eu/article/11/>

Please read the following three generated narratives. After this, you are asked to rank these narratives from worst to best, in terms of how well they represent compliance with Technical Documentation, as described in article 11 (1 best - 3 worst). You can find this question by scrolling down, as it is positioned right below the three narratives:

Zero-shot prompting:

Megan's team had almost finished their new AI tool for helping hospitals prioritise equipment maintenance, but before anyone talked about launching it, she opened a folder called "How we know this system is safe enough." Into it, she placed the story of the system: what it was built to do, which data had shaped it, how it had been tested, what risks they had found, which limits the hospital staff needed to know, and what safeguards were built in. Each time the engineers changed something important, Megan made them update the folder too, because a document that only described yesterday's system would be useless tomorrow. When the hospital's compliance officer visited, Megan did not try to explain everything from memory; she handed over the folder, clear enough for someone outside the team to understand how the AI met the legal requirements. Because the tool was part of a regulated medical product, she also made sure the same file covered both the AI Act information and the product-law information, instead of creating two confusing piles of paperwork. Her small start-up used the simplified documentation form meant for SMEs, but she treated it seriously: not as a shortcut around compliance, but as a clearer way to prove it. By launch day, the team understood that technical documentation was not decoration for auditors; it was the system's passport, prepared before the AI entered the real world, kept up to date, and detailed enough for authorities or notified bodies to check whether the high-risk AI system truly complied.

Contextual prompting:

Madison and John, two former UvA data science students, had built an AI tool that helped small clinics decide which patient files needed urgent review. At first, they thought their biggest job was making the model accurate, but their advisor reminded them that a system used in healthcare would not be judged only by what it could do on demo day. “Imagine an inspector walks in six months from now,” she said. “Could you show not just that it works, but why it was built this way, what data shaped it, what risks you expected, what tests you ran, and how a clinic is supposed to use it safely?” So before offering the tool to their first clinic, they opened a shared folder they called the system’s “passport.” In it, they described the purpose of the AI, the main design choices, the data used for training and testing, the performance results, the limits they had found, the risks they had identified, and the safeguards they had added, such as human review before any urgent-care recommendation was acted on. When they changed the model after discovering that it performed less well on older patient records, they did not treat the documentation as finished paperwork; they updated the passport so it still matched the real system. Because they were a small startup, they used the simplified form made for companies like theirs, but they still made sure the explanation was complete enough that an outside authority could follow the trail from idea, to development, to testing, to safe deployment. Later, when a clinic asked whether the tool was “EU AI Act ready,” Madison did not send a marketing slide. She sent the technical file. The clinic’s compliance officer could see that the team had not merely built an AI product and hoped for trust afterwards; they had kept a clear, living record that made the system inspectable, accountable, and ready to be assessed before people depended on it.

Rephrase and Respond prompting:

Before Lina's hospital could start using its new AI tool to help doctors spot early signs of illness, she refused to let the team simply plug it in and hope for the best. The software seemed impressive in testing, but Lina knew that trust could not depend on promises alone. She asked the engineers to gather everything that explained how the system worked: what it was designed to do, what data had been used to build and test it, what risks they had found, how they had reduced those risks, how accurate the system was, what limits doctors needed to know, and what should happen if something went wrong. At first, one developer complained that this felt like writing a manual for people who would never read it, but Lina explained that the file was more like the system's passport. Without it, no one outside the team could understand where the AI came from, what it was safe to use for, or whether it followed the required rules. A few weeks later, when the hospital changed the way the tool displayed warnings, Lina made sure the documentation changed too, because an old description of a new system would be useless. When an external assessor later asked how the hospital knew the AI was ready, Lina did not rely on a confident speech or a glossy demo. She opened the documentation and showed the whole trail: the design choices, the tests, the safeguards, the instructions, and the updates. The assessor could follow the story of the system from idea to deployment, and that was exactly the point: technical documentation made the AI understandable, checkable, and accountable before it touched real patients.

Zero-shot prompting

Contextual prompting

Rephrase and Respond prompting

Based on the one generated narrative that you just ranked as the best representation of Article 11 - Technical Documentation, please evaluate the following statements below:

	Strongly disagree	Somewhat disagree	Neutral	Somewhat agree	Strongly agree
The narrative contains correct information, avoiding any errors or inaccuracies compared to Art. 11	☐	☐	☐	☐	☐
The narrative is based on the Art. 11, excluding any extraneous or unnecessary information	☐	☐	☐	☐	☐
The narrative contains relevant information and represents Art. 11 well	☐	☐	☐	☐	☐

How well does the following multiple-choice question examine whether non-experts have understood the regulatory concept of Technical Documentation, as described in Article 11 of the AI Act?

Question: According to Article 11 of the EU AI Act, what must providers of high-risk AI systems do before placing the system on the market or putting it into service?

- (A) Prepare and keep updated technical documentation that clearly shows how the system meets legal requirements and allows authorities or assessors to check compliance.
- (B) Wait until the AI system has already been used in real-world settings before writing documentation, so the file reflects actual performance rather than planned safeguards.
- (C) Create only a short marketing summary explaining the system's

benefits, unless a regulator later asks for more technical information.

(D) Keep technical details private within the development team, because external authorities only need to know whether the provider promises compliance.

Not well ☒ at all

Slightly ☐ well

Moderately ☐ well

Very ☐ well

Extremely ☐ well

Do you have any further feedback on how your best ranked narrative and the multiple-choice question could be improved? (optional)

art 14

Regulatory Concept #4:

Below, you can find a link that takes you to *article 14* from the AI Act about *Human Oversight*. It is **not mandatory** to read this, but it is available if you would like to scan through it, for example, as this might help you answer the upcoming question.

<https://artificialintelligenceact.eu/article/14/>

Please read the following three generated narratives. After this, you are asked to rank these narratives from worst to best, in terms of how well they represent compliance with Human Oversight, as described in article 14 (1 best - 3 worst). You can find this question by scrolling down, as it is positioned right below the three narratives:

Zero-shot prompting:

At the hospital, Mira's team introduced an AI tool that helped rank urgent scan results, but they never treated it like the doctor in charge. Before using it, the supplier showed them exactly what the tool was built for, where it could make mistakes, and what warning signs to watch for. Mira and two senior nurses were trained to read its recommendations, question them, and notice when the system behaved oddly. One afternoon, the tool marked an elderly patient's scan as low priority, but Mira saw that the patient's symptoms did not fit the usual pattern. Remembering that people can become too trusting of automated suggestions, she paused, reviewed the scan herself, and asked a radiologist to check it too. They found a serious issue and overrode the AI's recommendation. Later, when the tool began giving several unusual rankings, Mira used the clinic's stop procedure and reported the problem instead of letting it continue. In practice, the clinic complied with human oversight not by having a person merely "near" the AI, but by making sure trained people could understand its limits, monitor it, challenge it, override it, and stop it safely whenever health, safety, or rights might be at risk.

Contextual prompting:

Amber and Noor had just launched their AI startup in Amsterdam, building a tool that helped hospitals prioritise patient appointment requests. At first, the system seemed almost magical: it read incoming forms, noticed urgent words, compared symptoms, and suggested which

patients should be seen first. But before they let a hospital use it, Noor added a dashboard that showed not only the system's recommendation, but also why it had made that suggestion, how confident it was, and whether the case looked unusual. Amber then trained the hospital coordinators to treat the AI like a very fast assistant, not like a final decision-maker. One coordinator, Sara, was given clear authority to question the system, override it, or pause it completely if something felt wrong. On the first week, the AI marked an elderly patient's request as "low priority" because the form was short and calm. Sara noticed the warning panel showing that the system had little information and that short descriptions could reduce accuracy. Instead of accepting the recommendation, she opened the patient's file, saw a history of heart problems, and moved the appointment up. Later that day, she reported the issue to Amber and Noor, who adjusted the instructions and added a stronger alert for cases with too little information. The AI had helped sort hundreds of requests, but it had not been allowed to run the process on its own. A person who understood its strengths and weaknesses stayed in charge, watched for mistakes, resisted the temptation to trust it automatically, and could step in before the system's recommendation harmed someone.

Rephrase and Respond prompting:

At a busy hospital, Sofia worked as a senior nurse in a department that had recently started using an AI system to help assess which patients might need urgent attention first. The system was fast and often useful, but the hospital never treated it like a replacement for medical judgment. Before Sofia and her colleagues were allowed to use it, they were trained on what the system could do, where it could make mistakes, and which signs might show that something was wrong with its recommendations. Each morning, Sofia checked the AI dashboard while also speaking with patients, reading their files, and listening to the

doctors. One afternoon, the system marked an elderly man as low priority, but Sofia noticed that his breathing looked unusual and that his symptoms did not fit the system’s ordinary pattern. Instead of simply accepting the AI’s output, she paused, reviewed the information, and called a doctor to examine him immediately. She also recorded why she had disagreed with the system’s recommendation. Later that day, when the AI began producing several strange results in a row, Sofia used the hospital’s procedure to suspend its use and alerted the technical team. The hospital did not see this as a failure of the AI project; it saw it as the system working as it should, with people remaining informed, attentive, and able to intervene. In practice, human oversight meant that the AI could support decisions, but it could not quietly take control of them.

Zero-shot prompting
Contextual prompting
Rephrase and Respond prompting

Based on the one generated narrative that you just ranked as the best representation of Article 14 - Human Oversight, please evaluate the following statements below:

	Strongly disagree	Somewhat disagree	Neutral	Somewhat agree	Strongly agree
The narrative contains correct information, avoiding any errors or inaccuracies compared to Art. 14	☐	☐	☐	☐	☐
The narrative is based on the Art. 14, excluding any extraneous or unnecessary information	☐	☐	☐	☐	☐

	Strongly disagree	Somewhat disagree	Neutral	Somewhat agree	Strongly agree
The narrative contains relevant information and represents Art. 14 well	☐	☐	☐	☐	☐

How well does the following multiple-choice question examine whether non-experts have understood the regulatory concept of Human Oversight, as described in Article 14 of the AI Act?

Question: In the context of Article 14 of the EU AI Act, what is the main requirement for using a high-risk AI system such as an AI tool that prioritises hospital patients?

- (A) The AI system may make final decisions as long as it is usually accurate and helps process cases faster.
- (B)** Human staff must be able to understand the AI system’s limits, monitor its outputs, question or override its recommendations, and stop its operation when needed.
- (C) The AI system only needs to explain its recommendations after a mistake has already caused harm.
- (D) Human oversight is satisfied if a person is physically present while the AI system operates, even if they cannot interpret or challenge its outputs.

Not ☒ at all	Slightly ☒ well	Moderately ☒ well	Very ☒ well	Extremely ☒ well
---	--	--	--	---

Do you have any further feedback on how your best ranked narrative and the multiple-choice question could be improved? (optional)



Powered by Qualtrics