

Default Question Block

Thank you for considering taking part in this study.

This survey is part of a Master's thesis project on leveraging Large Language Model to generate short narratives about regulatory concepts of the EU Artificial Intelligence Act (AI Act). The purpose of the generated narratives is to simplify the meaning of these regulatory concepts for non-experts and to clarify how this target group is expected to comply with such concepts.

A selection of 4 regulatory concepts is made, of which you are going to read two legal text summaries first, and representative generated narratives (200-400 words) after. After reading each text, you are asked to evaluate your reading experience based on a small amount of Likert scale questions, and your comprehension is going to be examined through a multiple-choice question.

This survey is estimated to take 10-15 minutes and offers a chance to receive a €50 Amazon giftcard or one of the two €25 giftcards as a token of appreciation for completing this survey.

By continuing, you confirm that you have read this information and agree that your responses may be collected. Your responses will be used for academic research purposes only.

Are you currently a student at the University of Amsterdam (UvA)?*

- ☐ Yes
- ☐ No, I already graduated from UvA
- ☐ No, I have never been

Block 1

Before we move on to the survey questions, you are asked to fill in a few questions about your background. This information helps accurately describe the participant sample and verify that your data meets the criteria for this user study.

What is your student email address? (required if you want a chance at receiving the €50 or one of the two €25 Amazon giftcards)

What course are you currently taking?*

- ☐ MSc Information Systems
- ☐ MSc Data Science

- ☐ MSc Artificial Intelligence
- ☐ Other, ...

Do you have legal expertise regarding the AI Act?*

- ☐ Yes
- ☐ No

Block 3

Regulatory concept #1 - Article 9 Risk Management System

Read the following legal text summary carefully, the upcoming questions will be about this text:

The EU AI Act requires a risk management system for high-risk AI systems. This system should be a continuous process throughout the AI's lifecycle, regularly reviewed and updated. It should identify and analyze potential risks to health, safety, or fundamental rights, estimate and evaluate these risks, and adopt measures to manage them. These measures should balance minimizing risks and fulfilling requirements. The system should also ensure that any remaining risk is acceptable and that measures are in place to eliminate or reduce risks as much as possible. High-risk AI systems should be tested to identify the best risk management measures and to ensure they work as intended. The system

should also consider whether the AI could negatively impact people under 18 or other vulnerable groups.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

Did the legal text summary contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this text:

Did the legal text summary contain sufficient information to understand this regulatory concept?*

1 (bad): the legal text summary does not contain rich and detailed

information

5 (good): the legal text summary contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this text:

How was your overall reading experience while reading this legal text summary?*

1 (bad): the legal text summary is highly difficult to read with rare words and complex structures

5 (good): the legal text summary is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the legal text summary that you have read, answer the following multiple-choice question (only one right answer):

In the EU AI Act's risk management system for high-risk AI systems, what is the main responsibility of the provider?*

- ☐ To prove before launch that the AI system is accurate enough, after which no further risk review is needed unless a serious incident occurs.
- ☐ To maintain a continuous process that identifies, evaluates, reduces, tests, documents, and regularly updates controls for risks to health, safety, and fundamental rights throughout the AI system's lifecycle.
- ☐ To transfer responsibility for all remaining risks to the human users by giving them warnings, training materials, and instructions before deployment.
- ☐ To eliminate every possible risk completely before the AI system is placed on the market, even if doing so makes the system unable to function.

narrative 2

Regulatory concept #2 - Article 10 Data Governance

Read the following generated legal text summary carefully, the upcoming questions will be about this text:

This article states that high-risk AI systems must be developed using high-quality data sets for training, validation, and testing. These data sets should be managed properly, considering factors like data collection processes, data preparation, potential biases, and data gaps. The data sets should be relevant, representative, error-free, and complete as much as possible. They should also consider the specific context in which the AI system will be used. In some cases, providers may process special categories of personal data to detect and correct biases, but they must follow strict conditions to protect individuals' rights and freedoms.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

Did the legal text summary contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this text:

Did the legal text summary contain sufficient information to understand this regulatory concept?*

1 (bad): the legal text summary does not contain rich and detailed information

5 (good): the legal text summary contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this text:

How was your overall reading experience while reading this legal text summary?*

1 (bad): the legal text summary is highly difficult to read with rare words and complex structures

5 (good): the legal text summary is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the legal text summary that you have read, answer the following multiple-choice question (only one right answer):

A startup is preparing a high-risk AI system for use in hospitals. Which approach best reflects the data governance requirements described in Article 10?*

- ☐ The startup uses a large hospital dataset because the model performs well in a demo, but it does not investigate whether the data reflects the patients and clinical settings where the system will actually be used.
- ☐ The startup checks where the data came from, why it was collected, how it was cleaned and labelled, whether important groups or settings are missing, whether bias could affect outcomes, and whether the training, validation, and testing data are suitable for the system's intended use.
- ☐ The startup focuses mainly on whether the AI system is accurate overall, because Article 10 is about model performance rather than the way datasets are collected, prepared, and governed.
- ☐ The startup avoids looking at sensitive group information under all circumstances, even when this would be strictly necessary and carefully protected to detect and correct bias.

legal 1

Regulatory concept #3 - Article 11 Technical documentation

Read the following generated narrative carefully, the upcoming questions will be about this story:

Madison and John, two former UvA data science students, had built an

AI tool that helped small clinics decide which patient files needed urgent review. At first, they thought their biggest job was making the model accurate, but their advisor reminded them that a system used in healthcare would not be judged only by what it could do on demo day. “Imagine an inspector walks in six months from now,” she said. “Could you show not just that it works, but why it was built this way, what data shaped it, what risks you expected, what tests you ran, and how a clinic is supposed to use it safely?” So before offering the tool to their first clinic, they opened a shared folder they called the system’s “passport.” In it, they described the purpose of the AI, the main design choices, the data used for training and testing, the performance results, the limits they had found, the risks they had identified, and the safeguards they had added, such as human review before any urgent-care recommendation was acted on. When they changed the model after discovering that it performed less well on older patient records, they did not treat the documentation as finished paperwork; they updated the passport so it still matched the real system. Because they were a small startup, they used the simplified form made for companies like theirs, but they still made sure the explanation was complete enough that an outside authority could follow the trail from idea, to development, to testing, to safe deployment. Later, when a clinic asked whether the tool was “EU AI Act ready,” Madison did not send a marketing slide. She sent the technical file. The clinic’s compliance officer could see that the team had not merely built an AI product and hoped for trust afterwards; they had kept a clear, living record that made the system inspectable, accountable, and ready to be assessed before people depended on it.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this generated narrative:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this generated narrative:

Did the narrative contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this narrative:

Did the generated narrative contain sufficient information to understand this regulatory concept?*

1 (bad): the narrative does not contain rich and detailed information

5 (good): the narrative contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this narrative:

How was your overall reading experience while reading this legal text summary?*

1 (bad): the narrative is highly difficult to read with rare words and complex structures

5 (good): the narrative is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the narrative that you have read, answer the following multiple-choice question (only one right answer):

According to Article 11 of the EU AI Act, what must providers of high-risk AI systems do before placing the system on the market or putting it into service?*

- ☐ Prepare and keep updated technical documentation that clearly shows how the system meets legal requirements and allows authorities or assessors to check compliance.
- ☐ Wait until the AI system has already been used in real-world settings before writing documentation, so the file reflects actual performance rather than planned safeguards.
- ☐ Create only a short marketing summary explaining the system's benefits, unless a regulator later asks for more technical information.
- ☐ Keep technical details private within the development team, because external authorities only need to know whether the provider promises compliance.

legal 2

Regulatory concept #4 - Article 14 Human Oversight

Read the following generated narrative carefully, the upcoming questions will be about this story:

At the hospital, Mira's team introduced an AI tool that helped rank urgent scan results, but they never treated it like the doctor in charge. Before using it, the supplier showed them exactly what the tool was built for, where it could make mistakes, and what warning signs to watch for. Mira and two senior nurses were trained to read its recommendations, question them, and notice when the system behaved oddly. One afternoon, the tool marked an elderly patient's scan as low priority, but Mira saw that the patient's symptoms did not fit the usual pattern. Remembering that people can become too trusting of automated suggestions, she paused, reviewed the scan herself, and asked a radiologist to check it too. They found a serious issue and overrode the AI's recommendation. Later, when the tool began giving several unusual rankings, Mira used the clinic's stop procedure and reported the problem instead of letting it continue. In practice, the clinic complied with human oversight not by having a person merely "near" the AI, but by making sure trained people could understand its limits, monitor it, challenge it, override it, and stop it safely whenever health, safety, or rights might be at risk, which reflects Article 14's requirements for high-risk AI systems.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this generated narrative:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this narrative:

Did the narrative contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this narrative:

Did the narrative contain sufficient information to understand this regulatory concept?*

1 (bad): the legal text summary does not contain rich and detailed information

5 (good): the legal text summary contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this narrative:

How was your overall reading experience while reading this narrative?*

1 (bad): the narrative is highly difficult to read with rare words and complex structures

5 (good): the narrative is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the narrative that you have read, answer the following multiple-choice question (only one right answer):

In the context of Article 14 of the EU AI Act, what is the main requirement for using a high-risk AI system such as an AI tool that prioritises hospital patients?*

- ☐ The AI system may make final decisions as long as it is usually accurate and helps process cases faster.
- ☐ Human staff must be able to understand the AI system's limits, monitor its outputs, question or override its recommendations, and stop its operation when needed.
- ☐ The AI system only needs to explain its recommendations after a mistake has already caused harm.
- ☐ Human oversight is satisfied if a person is physically present while the AI system operates, even if they cannot interpret or challenge its outputs.

Block 7

Thank you for participating in this survey!

Below you can find a link that takes you to an example website of how these regulatory concepts are currently displayed with their original text,

and with an added button for the narrative generation. Opening this link is completely optional. You may skip this and answer the final question below as well if you would like to.

<https://letter-bear-43346891.figma.site>

Please answer the final question on this statement to round this experiment up:

The previously presented generated narratives were useful to understand compliance with the given regulatory concepts*

Strongly ☒ disagree Somewhat ☒ disagree Neutral ☒ Somewhat ☒ agree Strongly ☒ agree

Block 2

Before we move on to the survey questions, you are asked to fill in a few questions about your background. This information helps accurately describe the participant sample and verify that your data meets the criteria for this user study.

What is your name?*

What is your email address? (required if you want a chance at receiving the €50 or one of the two €25 Amazon giftcards)

What course did you graduate from?*

- ☐ MSc Information Systems
- ☐ MSc Data Science
- ☐ MSc Artificial Intelligence
- ☐ Other, ...

Do you have legal expertise regarding the AI Act?*

- ☐ Yes
- ☐ No