

Default Question Block

Thank you for considering taking part in this study.

This survey is part of a Master's thesis project on leveraging Large Language Model to generate short narratives about regulatory concepts of the EU Artificial Intelligence Act (AI Act). The purpose of the generated narratives is to simplify the meaning of these regulatory concepts for non-experts and to clarify how this target group is expected to comply with such concepts.

A selection of 4 regulatory concepts is made, of which you are going to read two representative generated narratives (200-400 words) first, and two legal text summaries after. After reading each text, you are asked to evaluate your reading experience based on a small amount of Likert scale questions, and your comprehension is going to be examined through a multiple-choice question.

This survey is estimated to take 10-15 minutes and offers a chance to receive a €50 Amazon giftcard or one of the two €25 giftcards as a token of appreciation for completing this survey.

By continuing, you confirm that you have read this information and agree that your responses may be collected. Your responses will be used for academic research purposes only.

Are you currently a student at the University of Amsterdam (UvA)?*

- ☐ Yes
- ☐ No, I already graduated from UvA
- ☐ No, I have never been

Block 1

Before we move on to the survey questions, you are asked to fill in a few questions about your background. This information helps accurately describe the participant sample and verify that your data meets the criteria for this user study.

What is your student email address? (required if you want a chance at receiving the €50 or one of the two €25 Amazon giftcards)

What course are you currently taking?*

- ☐ MSc Information Systems
- ☐ MSc Data Science

- ☐ MSc Artificial Intelligence
- ☐ Other, ...

Do you have legal expertise regarding the AI Act?*

- ☐ Yes
- ☐ No

Block 3

Regulatory concept #1 - Article 9 Risk Management System

Read the following generated narrative carefully, the upcoming questions will be about this story:

Maya worked at a small company building an AI tool that helped hospitals decide which patients should be seen first in the emergency department. At first, the team was excited because the system seemed fast and accurate, but before anyone was allowed to use it, Maya gathered the doctors, nurses, engineers, and patient safety officer around one table and asked a simple question: “What could go wrong, even if everyone uses this exactly as intended?” They listed the obvious dangers first: the AI might underestimate a child’s symptoms, misunderstand rare conditions, perform worse for people whose data was underrepresented, or make staff trust its recommendation too much. Then Maya asked a second question: “How might people misuse it in a way we can reasonably predict?” The team imagined a crowded

night shift where a tired nurse might follow the AI score without checking the patient, or a hospital manager might use the tool outside the emergency department even though it was not designed for that. Instead of treating these worries as paperwork, Maya turned them into design decisions. The system was changed so that it could not make the final decision by itself, displayed warnings when cases were uncertain, explained the main reasons behind its recommendation, and clearly told staff when the tool should not be used. The team tested it again and again, including with difficult cases, children's cases, and examples from different patient groups, until the remaining risks were reduced as much as technically possible and judged acceptable. Every important choice was written down: the risks found, the tests run, the trade-offs made, the safeguards added, and the instructions hospitals would need. When the tool finally went live, Maya did not close the risk file and move on. She set up a routine to review incident reports, user feedback, and performance data from real hospitals, because new risks could appear after deployment. Months later, when the team noticed that the AI was less reliable in a certain type of paediatric case, they updated the system, revised the instructions, and recorded the change.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this narrative:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this narrative:

Did the narrative contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this narrative:

Did the narrative contain sufficient information to understand this regulatory concept?*

1 (bad): the narrative does not contain rich and detailed information

5 (good): the narrative contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this narrative:

How was your overall reading experience while reading this narrative?*

1 (bad): the story is highly difficult to read with rare words and complex structures

5 (good): the story is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the narrative that you have read, answer the following multiple-choice question (only one right answer):

In the EU AI Act's risk management system for high-risk AI systems, what is the main responsibility of the provider?*

- ☐ To prove before launch that the AI system is accurate enough, after which no further risk review is needed unless a serious incident occurs.
- ☐ To maintain a continuous process that identifies, evaluates, reduces, tests, documents, and regularly updates controls for risks to health, safety, and fundamental rights throughout the AI system's lifecycle.
- ☐ To transfer responsibility for all remaining risks to the human users by giving them warnings, training materials, and instructions before deployment.
- ☐ To eliminate every possible risk completely before the AI system is placed on the market, even if doing so makes the system unable to function.

narrative 2

Regulatory concept #2 - Article 10 Data Governance

Read the following generated narrative carefully, the upcoming questions will be about this story:

Mary and Jonas had just launched their Amsterdam AI startup, which built a tool to help hospitals decide which patients should be invited for an early diabetes screening. At first, their model looked impressive: it was fast, accurate in their demo, and investors loved it. But before they offered it to real clinics, Mary stopped the team and said, "We cannot just ask whether the model works. We need to ask whether the data

deserves to be trusted.” So they opened a shared document and traced every dataset back to its source: which hospital it came from, why it had originally been collected, whether patients had been informed properly, and whether the data was meant for this kind of prediction. Then they checked how the data had been cleaned, labelled, updated, and combined. One intern noticed that some older records used different names for the same diagnosis, while another found missing values in data from smaller clinics. Instead of ignoring those problems, the team wrote them down and decided how to fix or limit them. Jonas then compared the dataset with the people the system would actually be used for. He realised that most training data came from urban hospitals and did not include enough patients from rural areas, younger adults, or certain migrant communities. “If we train only on the people who are easiest to find,” he said, “the system may fail exactly when it meets people outside that pattern.” The team added more suitable data where they lawfully could, tested whether the model performed differently across groups, and looked for biases that could unfairly delay screening for some patients. When they needed sensitive information to check whether bias existed, they used it only where strictly necessary, protected access carefully, and deleted it once the bias check was done. By the time the startup met its first hospital client, Mary did not simply say, “Our AI is accurate.” She showed a record of the choices they had made: where the data came from, what it was supposed to represent, what gaps they found, how they reduced bias, and why the final training, validation, and testing data were suitable for the medical setting in which the tool would be used. The hospital’s compliance officer smiled and said, “So your data governance is not a folder you created at the end. It shaped the system before patients ever depended on it.”

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this narrative:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this narrative:

Did the narrative contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this narrative:

Did the narrative contain sufficient information to understand this regulatory concept?*

1 (bad): the narrative does not contain rich and detailed information

5 (good): the narrative contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this narrative:

How was your overall reading experience while reading this narrative?*

1 (bad): the story is highly difficult to read with rare words and complex structures

5 (good): the story is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the narrative that you have read, answer the following multiple-choice question (only one right answer):

A startup is preparing a high-risk AI system for use in hospitals. Which approach best reflects the data governance requirements described in Article 10?*

- ☐ The startup uses a large hospital dataset because the model performs well in a demo, but it does not investigate whether the data reflects the patients and clinical settings where the system will actually be used.
- ☐ The startup checks where the data came from, why it was collected, how it was cleaned and labelled, whether important groups or settings are missing, whether bias could affect outcomes, and whether the training, validation, and testing data are suitable for the system's intended use.
- ☐ The startup focuses mainly on whether the AI system is accurate overall, because Article 10 is about model performance rather than the way datasets are collected, prepared, and governed.
- ☐ The startup avoids looking at sensitive group information under all circumstances, even when this would be strictly necessary and carefully protected to detect and correct bias.

legal 1

Regulatory concept #3 - Article 11 Technical documentation

Read the following legal text summary carefully, the upcoming questions will be about this text:

This EU law states that before a high-risk AI system is launched, detailed technical documentation must be prepared and kept updated. This documentation should prove that the AI system meets the law's requirements and provide clear information for authorities to check compliance. It should include certain key elements. Small businesses, including start-ups, can provide this information in a simpler way. The EU will create a simplified form for this purpose. If a high-risk AI system is related to a product covered by other EU laws, one set of documentation should include all necessary information. The EU can update these requirements as technology advances.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

Did the legal text summary contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this text:

Did the legal text summary contain sufficient information to understand this regulatory concept?*

1 (bad): the legal text summary does not contain rich and detailed information

5 (good): the legal text summary contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this text:

How was your overall reading experience while reading this legal text summary?*

1 (bad): the legal text summary is highly difficult to read with rare words and complex structures

5 (good): the legal text summary is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the legal text summary that you have read, answer the following multiple-choice question (only one right answer):

According to Article 11 of the EU AI Act, what must providers of high-risk AI systems do before placing the system on the market or putting it into service?*

- ☐ Prepare and keep updated technical documentation that clearly shows how the system meets legal requirements and allows authorities or assessors to check compliance.
- ☐ Wait until the AI system has already been used in real-world settings before writing documentation, so the file reflects actual performance rather than planned safeguards.
- ☐ Create only a short marketing summary explaining the system's benefits, unless a regulator later asks for more technical information.
- ☐ Keep technical details private within the development team, because external authorities only need to know whether the provider promises compliance.

legal 2

Regulatory concept #4 - Article 14 Human Oversight

Read the following legal text summary carefully, the upcoming questions will be about this text:

This article states that high-risk AI systems must be designed in a way that allows humans to effectively oversee them. The goal of human oversight is to prevent or minimize risks to health, safety, or fundamental rights that may arise from using these systems. The oversight measures should match the risks and context of the AI system's use. These measures could be built into the system by the provider or implemented by the user. The AI system should be provided in a way that allows the overseer to understand its capabilities and limitations,

detect and address issues, avoid over-reliance on the system, interpret its output, decide not to use it, or stop its operation. For certain high-risk AI systems, any action or decision based on the system's identification must be verified by at least two competent individuals.

Please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

How clear is the regulatory concept to you now?*

1 (bad): still unclear or hard to understand

5 (good): completely clear or easy to understand

①

②

③

④

⑤

Again, please rate your perceived **clarity** of compliance with the represented regulatory concept after you have read this text:

Did the legal text summary contain confusing or distracting elements?*

1 (bad): there were many confusing/distracting elements

5 (good): there were no confusing/distracting elements

①

②

③

④

⑤

Please rate your perceived **informativeness** of this text:

Did the legal text summary contain sufficient information to understand

this regulatory concept?*

1 (bad): the legal text summary does not contain rich and detailed information

5 (good): the legal text summary contains rich and detailed information

①

②

③

④

⑤

Please rate your perceived **readability** of this text:

How was your overall reading experience while reading this legal text summary?*

1 (bad): the legal text summary is highly difficult to read with rare words and complex structures

5 (good): the legal text summary is easy to read, well-structured, and flows naturally

①

②

③

④

⑤

Based on the legal text summary that you have read, answer the following multiple-choice question (only one right answer):

In the context of Article 14 of the EU AI Act, what is the main requirement for using a high-risk AI system such as an AI tool that prioritises hospital patients?*

- ☐ The AI system may make final decisions as long as it is usually accurate and helps process cases faster.
- ☐ Human staff must be able to understand the AI system's limits, monitor its outputs, question or override its recommendations, and stop its operation when needed.

- ☐ The AI system only needs to explain its recommendations after a mistake has already caused harm.
- ☐ Human oversight is satisfied if a person is physically present while the AI system operates, even if they cannot interpret or challenge its outputs.

Block 7

Thank you for participating in this survey!

Below you can find a link that takes you to an example website of how these regulatory concepts are currently displayed with their original text, and with an added button for the narrative generation. Opening this link is completely optional. You can skip this and simply answer the final question below as well if you would like to.

<https://letter-bear-43346891.figma.site>

Please answer the final question on this statement to round this experiment up:

The previously presented generated narratives were useful to understand compliance with the given regulatory concepts*

Strongly disagree Somewhat disagree Neutral Somewhat agree Strongly agree

Block 2

Before we move on to the survey questions, you are asked to fill in a few questions about your background. This information helps accurately describe the participant sample and verify that your data meets the criteria for this user study.

What is your name?*

What is your email address? (required if you want a chance at receiving the €50 or one of the two €25 Amazon giftcards)

What course did you graduate from?*

- ☐ MSc Information Systems
- ☐ MSc Data Science
- ☐ MSc Artificial Intelligence
- ☐ Other, ...

Do you have legal expertise regarding the AI Act?*

- ☐ Yes
- ☐ No